

Project 5, FYS 3150 / 4150, fall 2013

Student #

December 8, 2013

1 Introduction

Open clusters are groups of up to a few thousands stars, held together by mutual gravitational attraction. Open clusters generally last for a few hundred millions years. It is also important to underline that the parameters of the cluster's stars are kind of constant, since they are made from the same material. We want to build a model for such a cluster, and study its "cold collapse" (where "cold" comes from the fact that the particles have no initial velocity). To do so, we will begin with the approximation that stars are particles, point masses. Then, we will use a simple smoothing algorithm to increase the numerical stability of our system. We have outlined a method to find an optimal smoothing parameter if one begins from a well-tested value. This value can then be scaled to fit other simulations with a surprisingly simple formula.

Numerous algorithms could have been used to simulate the behavior of the cluster. We chose to focus on the Leapfrog and on the fourth-order Runge Kutta methods. Each one of the numeric methods has its pros and cons; we seek to find the one which gives priority to the stability instead of the short-term accuracy, as well as a reasonable running time. Due to the number of particles in our cluster, we will be interested in the statistical properties of our system instead of in the specific characteristics of each particle.

Our goal in this report is to simulate and to study the behavior of a cluster, after having paid attention to the algorithm used to derive the results. The first part of the report will focus on the analytical properties of the model (including algorithms). We then proceed to give an overview of the code structure, and finally present and discuss our results.

2 Theory

2.1 Derivation of the expression for τ_{crunch}

We start with the parametric form of the Friedmann equations for a closed universe containing only dust in addition to curvature[1]:

$$R(\psi) = a(\psi)R_0 = \frac{R_0\Omega_{m0}}{2(\Omega_{m0} - 1)}(1 - \cos \psi) \quad (1)$$

where $a(\psi)$ is the dimensionless scale factor and

$$t(\psi) = \frac{\Omega_{m0}}{2H_0(\Omega_{m0} - 1)^{3/2}}(\psi - \sin \psi). \quad (2)$$

We will use $t = 0$ at the Big Bang singularity for our sub-universe. From these results we see that $t_{max} = t(\psi = \pi)$ and $t_{crunch} = t(\psi = 2\pi)$, thus the elapsed time between these events is

$$\tau_{crunch} = t_{crunch} - t_{max} = \frac{\pi\Omega_{m0}}{2H_0(\Omega_{m0} - 1)^{3/2}}.$$

The mass parameter is defined by

$$\Omega_{m0} = \frac{8\pi G\rho_0}{3H_0^2}$$

and for readability we will make the substitution

$$u^2 = 8\pi G\rho_0 - 3H_0^2$$

thus

$$(\Omega_{m0} - 1)^{-3/2} = \left(\frac{u^2}{3H_0^2}\right)^{-3/2} = 3\sqrt{3}H_0^3u^{-3}$$

$$\frac{\Omega_{m0}}{(\Omega_{m0} - 1)^{3/2}} = \frac{8\pi G\rho_0}{3H_0^2} 3\sqrt{3}H_0^3u^{-3} = 8\sqrt{3}\pi G\rho_0H_0u^{-3}$$

$$\tau_{crunch} = \frac{\pi}{2H_0} 8\sqrt{3}\pi G\rho_0H_0u^{-3} = 4\sqrt{3}\pi^2 G\rho_0u^{-3}$$

Now we remember that at the time when $\rho = \rho_0$, everything is at rest, so we have $H_0 = (\frac{\dot{a}}{a})_{\tau=0} = 0$. Inserting this, we get $u^2 = 8\pi G\rho_0$, and

$$\tau_{crunch} = 4\sqrt{3}\pi^2 G\rho_0 (8\pi G\rho_0)^{-3/2} = \sqrt{\frac{4^2 3\pi^4 G^2 \rho_0^2}{8^3 \pi^3 G^3 \rho_0^3}} = \sqrt{\frac{3\pi}{32G\rho_0}} \quad (3)$$

which is what we wanted to show.

2.2 Lack of a singularity in our model

The reason we do not see a singularity in our model is that we have assumed pressureless (i.e. collisionless) matter that is homogeneously distributed in the Friedmann equations. In our simulation we will see internal “pressure” as kinetic energy from the collapse is turned into random motions by near-collisions between the point mass particles. These motions will halt the collapse, causing the gravitationally bound particles to form more or less stable orbits around the center of mass.[2, 3] Keep in mind that the kinetic energy is not evenly distributed, so occasionally particles that receive more than their fair share will become unbound and may escape from the system before they have time to lose their energy. It is predicted that this should happen to a certain percentage of the particles, depending on N , with a lower bound of 50%.[4]

We say that the system is stable when it satisfies the virial theorem:

$$2\langle K \rangle = -\langle P \rangle$$

where $\langle P \rangle$ and $\langle K \rangle$ are the averages of the potential and kinetic energies (actually, these are supposed to be time averages for the total energies of the entire system, but it turns out that the averages of a point in time of the energies per particle is a good approximation to this[5]. It can be shown that this happens at time $\tau_{vir} = 0.81\tau_{crunch}$ when the sphere has collapsed to half its initial size, so we see that τ_{crunch} is a natural time scale for virialization to occur (some sources do in fact use τ_{crunch} to mark the point when the system is virialized).[2]

2.3 G in units of ly , M_{\star} and τ_{crunch}

With τ_{crunch} given in years, we can rewrite equation 3 as

$$G_{yr} = \frac{3\pi}{32\tau_{crunch}^2 \rho_0}.$$

Switching time units to τ_{crunch} , we get that $\tau_{crunch} = 1$ in these units, hence

$$G = \frac{3\pi}{32\rho_0} = \frac{\pi^2 R_0^3}{8\mu N} \quad (4)$$

where for the latter equality we have used the definitions of average mass $\mu = \frac{M}{N}$ and initial mass density for a sphere $\rho_0 = \frac{M}{V_0} = \frac{\mu N}{V_n(R_0)}$, where $V_n(R_0)$ is the volume of the n-dimensional ball with radius R_0 .

This means our gravitational constant and our time unit both depend on N , R_0 and μ . We can verify that the units for G are now correct with R_0 given in light years and μ given in solar masses.

2.4 Calculating ϵ automatically

The challenge with the ϵ values is that we want it to be as small as possible (to give more realistic results), yet make the number of ejected particles as small as possible, conserving as much of the total energy as we can. We chose to determine a good fit for ϵ experimentally (see figures 10 and 11), and then we determined how this value would scale to different simulations in the following way:

The basic premise is that when $r = \epsilon$ (i.e. when the ϵ term starts dominating the gravitational potential), we want to set an upper limit on how much a particle's velocity can change during one time step due to the gravitational attraction from one other particle. We want this upper limit to be invariant across simulations:

$$\Delta v_1 = \Delta v_2$$

$$a_1 \Delta t_1 = a_2 \Delta t_2$$

We will assume all particles have the average mass μ and set $r = \epsilon$:

$$\frac{G_1 \mu_1 \Delta t_1}{2\epsilon_1^2} = \frac{G_2 \mu_2 \Delta t_2}{2\epsilon_2^2}$$

Now, inserting the value of G from 4 with initial radii R_1 and R_2 , we get:

$$\frac{\pi^2 R_1^3 \mu_1 \Delta t_1}{2\epsilon_1^2 \cdot 8\mu_1 N_1} = \frac{\pi^2 R_2^3 \mu_2 \Delta t_2}{2\epsilon_2^2 \cdot 8\mu_2 N_2}$$

$$\frac{R_1^3 \Delta t_1}{\epsilon_1^2 N_1} = \frac{R_2^3 \Delta t_2}{\epsilon_2^2 N_2}$$

giving us this handy formula for ϵ_2 if we have a good match for ϵ_1 :

$$\epsilon_2 = \sqrt{\frac{N_1}{N_2} \left(\frac{R_2}{R_1} \right)^3 \frac{\Delta t_2}{\Delta t_1}} \cdot \epsilon_1 \quad (5)$$

2.5 Gravitational potential with modified gravity

Starting with the magnitude of the force

$$F = \frac{GMm}{r^2 + \epsilon^2} = \frac{GMm}{\epsilon^2} \frac{1}{\left(\frac{r}{\epsilon}\right)^2 + 1} = \frac{GMm}{\epsilon^2} \frac{1}{u^2 + 1}$$

where we have used the substitution $u = \frac{r}{\epsilon}$ which gives $\frac{du}{dr} = \frac{1}{\epsilon}$, hence $dr = \epsilon \cdot du$, and so

$$E_p = m\Phi = \int F dr = \frac{GMm}{\epsilon^2} \int \frac{1}{u^2 + 1} \epsilon du = \frac{GMm}{\epsilon} (\arctan(u) + C)$$

We want $E_p \rightarrow 0$ as $r \rightarrow \infty$, and since $\arctan(u) \rightarrow \frac{\pi}{2}$ as $u \rightarrow \infty$, we achieve this by choosing $C = -\frac{\pi}{2}$:

$$E_p = \frac{GMm}{\epsilon} \left(\arctan\left(\frac{r}{\epsilon}\right) - \frac{\pi}{2} \right) \quad (\epsilon > 0) \quad (6)$$

2.6 Volume of the n-ball

To calculate the gravitational constant with τ_{crunch} as the time unit in any dimension, we need the volume of the sphere in n dimensions to calculate the initial mean density ρ_0 . This is accomplished by the following formula[6]:

$$V_n = \frac{\pi^{\frac{n}{2}}}{\Gamma(\frac{n}{2} + 1)} \quad (7)$$

2.7 Uniform distribution in the n-ball

Our dimension-independent algorithm for generating uniformly distributed points inside the n-ball is[7]:

1. Generate random points on the surface of the unit n-ball (i.e. randomize the directions of the unit vectors):
 - (a) Generate normally distributed n-dimensional vector $\mathbf{x}_n = [x_1, x_2, \dots, x_n]$ where the normal distribution has $\mu = 0$ and $\sigma = 1$.
 - (b) Calculate the n-dimensional norm of the vector $|\mathbf{x}_n|$. We chose to let Armadillo handle this, but another easy way to do it is using the n-dimensional dot product and taking the square root of this: $|\mathbf{x}_n| = \sqrt{\mathbf{x}_n \cdot \mathbf{x}_n} = \sqrt{\sum_i x_i^2}$.

- (c) Turn it into a unit vector: $\mathbf{u}_n = \frac{\mathbf{x}_n}{|\mathbf{x}_n|}$. We refer to (source) for the proof that this is uniformly distributed in terms of direction.
- 2. Generate a radius that results in a uniform distribution within maximum radius R_0 :
 - (a) Generate a uniformly distributed value $r \in [0, 1]$.
 - (b) The desired radius that takes into account that the outer spherical shells have a larger surface area is $R = \sqrt[n]{r} \cdot R_0$, where n is the number of dimensions.

2.8 Algorithms

For both of our algorithms, the step interval looks like $[t_0; t_{max}]$. We will work on n_{Steps} . Let's define h as $h = \frac{t_{max} - t_0}{n_{Steps} - 1}$.

2.8.1 Runge-Kutta

We have $t_i = t_0 + i \cdot h$ for each i belonging to $[0, n_{Steps}]$. We can now define $x_i = x(t_i)$.

The general formulae of Simpson's rule is: $\int_{t_i}^{t_i+1} f(t, x) dt = \frac{1}{6} h (f(t_i, x_i) + 2f(t_{i+\frac{1}{2}}, x_{i+\frac{1}{2}}) + f(t_{i+1}, x_{i+1}))$.

Runge-Kutta's method defines four quantities, to predict and correct the value of x_{i+1} :

$$k_1 = f(t_i, x_i)$$

$$k_2 = f(t_{i+\frac{1}{2}}, x_{i+\frac{1}{2}}) \quad x_{i+\frac{1}{2}} = x_i + \frac{h}{2} k_1$$

$$k_3 = f(t_{i+\frac{1}{2}}, x_{i+\frac{1}{2}}^*) \quad y_{i+\frac{1}{2}} = x_i + \frac{h}{2} k_2$$

$$k_4 = f(t_{i+1}, x_{i+1}) \quad x_{i+1} = x_i + h k_3$$

In our case, we will have to compute this results six times by time step: once for the x, y and z positions, and once for the x, y and z velocities.

This Runge-Kutta algorithm is especially interesting since we only need initial conditions to unfold it. Its approximation error runs like $\mathcal{O}(\Delta h^2)$.

2.8.2 Leapfrog algorithm

The Leapfrog algorithm is given by the three following steps:

$$x^{(1)}(t + \frac{h}{2}) = (x^{(1)}(t) + \frac{h}{2} * x^{(2)}(t) + \mathcal{O}(\Delta h^2))$$

$$x(t + h) = x(t) + x^{(1)}(t + \frac{h}{2}) + \mathcal{O}(\Delta h^3)$$

$$x^{(1)}(t + h) = x^{(1)}(t + \frac{h}{2}) + \frac{h}{2} * x^{(2)}(t + h) + \mathcal{O}(\Delta h^2)$$

Now, let's take a look at the approximation error for this algorithm. The first step to derive this algorithm is to use Taylor expansion on the position:

$$x(t+h) = x(t) + hx^{(1)}(t) + \frac{h^2}{2}x^{(2)}(t) + \mathcal{O}(\Delta h^3) = x(t) + h(x^{(1)}(t) + \frac{h}{2}x^{(2)}(t)) + \mathcal{O}(\Delta h^3)$$

In the same way, we have for the velocity:

$$x^{(1)}(t + \frac{h}{2}) = x^{(1)}(t) + \frac{h}{2} * x^{(2)}(t) + \mathcal{O}(\Delta h^2)$$

$$x^{(1)}(t - \frac{h}{2}) = x^{(1)}(t) - \frac{h}{2} * x^{(2)}(t) + \mathcal{O}(\Delta h^2)$$

$$x^{(1)}(t + \frac{h}{2}) - x^{(1)}(t - \frac{h}{2}) = x^{(1)}(t) + \frac{h}{2} * x^{(2)}(t) + \mathcal{O}(\Delta h^2) - (x^{(1)}(t) - \frac{h}{2} * x^{(2)}(t) + \mathcal{O}(\Delta h^2)) = h * x^{(2)}(t) + \mathcal{O}(\Delta h^2)$$

So

$$x^{(1)}(t + h) = x^{(1)}(t + \frac{h}{2}) + \frac{h}{2} * x^{(2)}(t + h) + \mathcal{O}(\Delta h^2)$$

We can rewrite the first step thanks to the expression of $x^{(1)}(t + \frac{h}{2})$:

$$x(t + h) = x(t) + x^{(1)}(t + \frac{h}{2}) + \mathcal{O}(\Delta h^3)$$

In the end, we can see that the approximation error runs like $\mathcal{O}(\Delta h^2)$ for the Leapfrog algorithm.

3 Implementation

3.1 Overview

The code is designed to be highly modular and independent of dimension to maximise reusability.



Figure 1: Class Diagram

3.2 Main program (Project5)

Sets up a series of simulations so we can plot certain results as a function of either N (number of particles), ϵ (correction factor for gravitation) or Δt (time step). This automated approach made the process of plotting considerably easier for us than setting up everything “by hand” each time. The main program also contains code to test the simulation against a 2D benchmark (Project 3) and contains helper methods for curve fitting.

3.3 SolarSystem class

Container class for the entire N-body system. This class is dimension independent (number of dimensions is given as a parameter), to make it as general and reusable as possible. Does not specify a method to solve the equations of motion, so any numerical method can be used on this class to update the positions every time step (by iterating over the celestial bodies contained in this class). A SolarSystem can make deep copies of itself (and all its CelestialBody and Gravity objects), which we make use of to be able to run different algorithms on identical copies of a system (for comparison of results and stability analysis). Some resource intensive properties, like potential energy, are only calculated when we need to plot the data - this can be set not to happen at every step to speed up the execution.

3.4 CelestialBody class

Particle class of our N-body simulation, with position and mass. Handles all calculations on the individual particle level (e.g. kinetic energy) and can be set to fixed if desired. Also stores properties calculated by the encompassing SolarSystem (like potential energy) so these will not have to be calculated on the fly every time they are needed. This class inherits dimensionality from the SolarSystem it belongs to, so we avoid creating a sub-dimensional celestial body by accident.

3.5 Gravity class

This class allows each SolarSystem to use a different gravity. One can change both the value of the gravitational constant to fit the time unit of choice and set an ϵ value to dampen collisions. For future use it is possible to extend this class to include any form of modified gravity (which is an object of interest in cosmology[8]). The only thing the SolarSystem class requires of this object is that it is able to return a force and potential energy when two CelestialBody objects are given as input. This way, a SolarSystem does not need to handle the specifics of gravitational forces itself, making the code more modular.

3.6 CelestialBodyInitializer class

Sets up a uniform position distribution given an initial max radius (generalized for any number of dimensions). Also generates normal distributed random masses. Doing this in a separate class allows the SolarSystem class to be as general as possible (we might not want a random distribution every time, but when we do, this class will provide it). Keeping this class dimension independent makes it more general and easier to use in other projects.

3.7 Solvers class

Contains code to iterate over a SolarSystem object over a number of time steps, using the Leapfrog, Runge-Kutta (4th order) and Euler-Cromer algorithms to solve the equations of motion. Creates copies of the system given as a parameter to do this, so we can keep the original system unchanged and also solve using several numerical methods simultaneously and compare the results. Keeping this in a separate class allows adaptation of this code to use other objects than the SolarSystems we employ here.

3.8 GaussPDF class

This is simply a static class wrapper for the code given at:

<https://www.uio.no/studier/emner/matnat/fys/FYS3150/h13/gaussiandeviate.cpp>

It provides random numbers in the uniform and normal probability distributions.

4 Results and analysis

4.1 Benchmarks and validation

We worked during a previous project on the implementation of the Runge Kutta algorithm to simulate the behavior of the Solar System. To test our code, we also wrote the Euler-Cromer algorithm. The outputs of our previous project were validated. In order to testify that our code works, we tested it against our previous results, in two dimensions, then we extended it in three dimensions.




Figure 2: Simulation with our Leapfrog algorithm




Figure 3: Simulation with our Runge Kutta - 4 algorithm




Figure 4: Simulation of our previous project

We found equivalent results between the two projects, for the same time step size, one point per day, and the same length of simulation, two years, which lead us to say that our current code is working.

When we use them for very long times, we can see that the Leapfrog algorithm is way more well-behaved than the Runge Kutta one. Thus, we can say that the Runge Kutta algorithm is not as stable as the Leapfrog one for when being processed for long times


We saw in the theory part that the forth-order Runge Kutta has an approximation error smaller than the Leapfrog algorithm. It basically means that the local data about each particles will be more accurate but since the Leapfrog method is more stable than the Runge Kutta method, and gives satisfying results for

small enough time steps. The Runge Kutta method is more accurate to compute the velocity, but here, it does not matter so much. It is also important to underline the fact that the Runge Kutta method generates an oscillation on the energy, while the energy is stable with the Leapfrog one. And if the time step gets to big, for both of our algorithm, everything becomes messy, and not usable.

Execution time (s)	For 1000 steps, two bodies	For 1000 steps, 100 bodies
Leap-frog	19,18	86,35
Forth-order Runge Kutta	30,85	171,25

Table 1: Elapsed Time to process the two body system

We can see here that the Runge Kutta algorithm takes more time to process the same set of data than the Leapfrog. For small data set, the time spent by the forth order Runge Kutta algorithm is twice as much as the time spent by the Leapfrog method. And when we increase the number of steps, or the number of bodies, we increase the time spent during processing. Thus, the difference in the execution time between Runge Kutta and Leapfrog becomes more and more significant.




```

E.k before: 1.33206e+033
E.p before: 0
E.tot before: 1.33206e+033
Entering the Benchmark part
Using Leap-frog
aaand ... We're done
E.tot after: -1.62099e+033
E.k after: 1.37681e+033
E.p after: -2.9979e+033

```

Figure 6:
Conservation of
the Energy –
Leap-frog

Figure 5: Conservation of the Energy – Leap-frog – Length = 2 years ~



```

E.k before: 1.33206e+033
E.p before: 0
E.tot before: 1.33206e+033
Entering the Benchmark part
Using Leap-frog
aaand ... We're done
E.tot after: -1.62099e+033
E.k after: 1.37681e+033
E.p after: -2.9979e+033

```

Figure 8:
Conservation
of the energy
– forth-order
Runge Kutta –
A changer !!!

Figure 7: Conservation of the energy – forth-order Runge Kutta = Length = 2 years ~

4.2 Application to a multi-bodies system

4.2.1 Evaluation of G

```
T_CRUNCH = 7.97246e+006  
G = 9.92352  
G_YLS = 1.56128e-013
```

Figure 9: Computation of G

With a system of $N = 100$ particles, an initial radius of $R_0 = 20$ light years, we found a τ_{crunch} of nearly 8 millions of years. In the figure above, the variable called G_YLS is the gravitational constant, uses meters, seconds, and kg. G is computed in ly, M_{\odot} and τ_{crunch} .

τ_{crunch} is evaluated in millions of years, thus, its unit could be the megannum (i.e. the unit for 1,000,000 years).

4.2.2 Choice of the algorithm

We will look for a satisfying time step. To do so, we will look at the total energy conservation over different time steps, for our two algorithms. Our system is the following: 100 particles, and an initial radius of 20 light years.

According to the time required to process equivalent results, the Leapfrog algorithm seems more interesting. Furthermore, we are going now to process on a N -body system, where N is at least equal to 100. Thus, we will be more interested in the statistical properties of our system, than in the local accuracy of the position or the velocity.

4.3 Energy conservation

4.3.1 Bound particles

When we refer to bound particles, we refer to those particles that are not gravitationally bound. Since potential energy is negative (at most 0), and kinetic energy is positive (at least 0), we say that a particle is bound when its total energy $E_{tot} = E_k + E_p < 0$.

4.3.2 Finding a reference value for ϵ

With a system of $N = 100$ particles, an initial radius of $R_0 = 20$ light years and a time step of $\frac{1}{250}\tau_{crunch}$, we did a series of simulations where we varied the gravitational correction parameter in the interval $\epsilon \in [0, 0.15]$. According to

equation 5, the values are supposed to be independent of mass, but the masses used were a normal distribution with mean $\mu = 10$ and $\sigma = 1$ (both in solar masses).

The criteria for accepting an ϵ value are that the number of bound particles are as high as possible, while the total energy, at least for the bound particles, is as conserved as possible. In the process, we discovered that these criteria are actually mutually exclusive (see figure 10): The energy conservation for bound particles improves if we eject many particles. Our interpretation of this result is that particles that are bound but close to being unbound will impact the energy bound conservation negatively: If these particles are actually ejected, the remaining particles' total energy is better conserved. Thus, we need to look at both criteria simultaneously when deciding on a value for ϵ . Additionally, we want ϵ to be as close to 0 (and the Newtonian limit) as possible, so we get more realistic results.

We also see that the number of ejected particles will fluctuate considerably between simulations with the same starting parameters.[4] Ideally, we would run many simulations with the same epsilon value to eliminate statistical errors - this is something that could be looked at in a future study.




Figure 10: Fraction of bound particles at end of simulation for a range of ϵ values.

Figure 11: Energy conservation (relative to starting potential energy) for a range of ϵ values.

The relative energy conservation in figure 11 was calculated using the following formula, using the total energy before the simulation as our scale:

$$\Delta E_{rel} = \frac{E_{after} - E_{before}}{|E_{before}|}$$

where E_{after} can be the total energy of all the particles and the bound particles, respectively.

The results show that the bound particle fraction reaches a value close to the extrapolated value for $N = 100$ (see fig. 3 in [4]) with as low a value as $\epsilon \approx 0.04$. The energy conservation takes a higher value to approach desired levels. Around $\epsilon \approx 0.10$ we notice that the two curves start behaving opposite of one another: When one value of ΔE_{rel} goes up, the other goes down for the same ϵ . Note also that the curve for the bound particles starts mirroring the fraction in figure 10. This is a result of the low number of particles: We expect the bound particle fraction to fluctuate considerably between simulations with the same starting parameters (again, see fig. 3 in [4]). Thus, when the energy difference for bound particles starts mirroring the number of bound particles, we interpret this as the curve having stabilized (the relative energy difference for all the particles seems to reach a plateau at this point). On the other hand, when $\epsilon < 0.10$, we do not see this behaviour: Both curves are in decline. Hence our (admittedly shaky) basis for choosing $\epsilon = 0.10$ is that after this point, we can no longer distinguish improvements in energy conservation from random fluctuations caused by variations in the fraction of bound particles.

Using this result, we propose inserting the following values into equation 5: $\epsilon_1 = 0.1$, $\Delta t_1 = \frac{1}{250}$, $N_1 = 100$, $R_1 = 20$. We then get the following formula that we have used in the remaining simulations for automating the process of choosing an acceptable value for ϵ :

$$\epsilon = \sqrt{\frac{100}{N} \left(\frac{R_0}{20}\right)^3 \frac{\Delta t}{\frac{1}{250}}} \cdot 0.1 = \frac{1}{4} \sqrt{\frac{R_0^3 \Delta t}{2N}} \quad (8)$$

which we expect could be of use in future versions of this project. This formula gave us consistently acceptable results with regards to energy conservation and particle rejection. However, for future projects, we would look at running many more simulations like the one we used to pick the reference value $\epsilon = 0.1$, to ensure that the process is statistically sound (and preferably with a higher number of particles): While the value we chose has been good enough in practice, we do not claim to have found the optimal value for the case we were studying, and expect future projects to be able to achieve much more accurate results using the method we have outlined above. Another interesting prospect is to seek a closed form solution for a reference value for ϵ (before scaling it), instead of taking the brute force approach and iterate over a set of possible values like we did here.

One should also keep in mind that the way we scale ϵ is linear in the sense that we only consider one pair of particle at a time, not looking at “crosstalk” from multi-particle collisions. Furthermore, we have assumed that all the particles have the same mass: If the standard deviation is significantly greater than the 10% of the mean we used to arrive at this result, the above formula may well not be applicable. That said, the particles that receive a change in velocity

that is too high for a timestep would tend to be the lighter particles, so we would expect to see the fraction of bound particles decrease while the energy conservation should feel the impact of a greater mass variance less in this case.

4.4 Particles ejection and equilibrium

As discussed in the theory part, the system seems to reach an equilibrium. We can see it happens around $1\tau_{crunch}$. And we know, after deriving the theoretical results, that it, in fact, happens after $0.81\tau_{crunch}$ (i.e. for circa 6,458 millions of years).

After this time, we can this on the following plots displaying the number of bound particles in regard to the time, that suddenly, the number of ejected particles blows up. And we can see, as expected, that without using the smoothing function, we lose a lot more of particles in just a few τ_{crunch} . The results are stabilized when we add epsilon, and the quantity of lost particles decreased.

When we increase N, the number of initial particles, we increase the percentage of lost particles too. But we keep the total mass constant, which means that as N increases, the particles become lighter. Since the energies, kinetic and potential, are both a function of the mass of the processed particles, we can see that the quantity of lost energy increases sharply with N and is nearly linear:

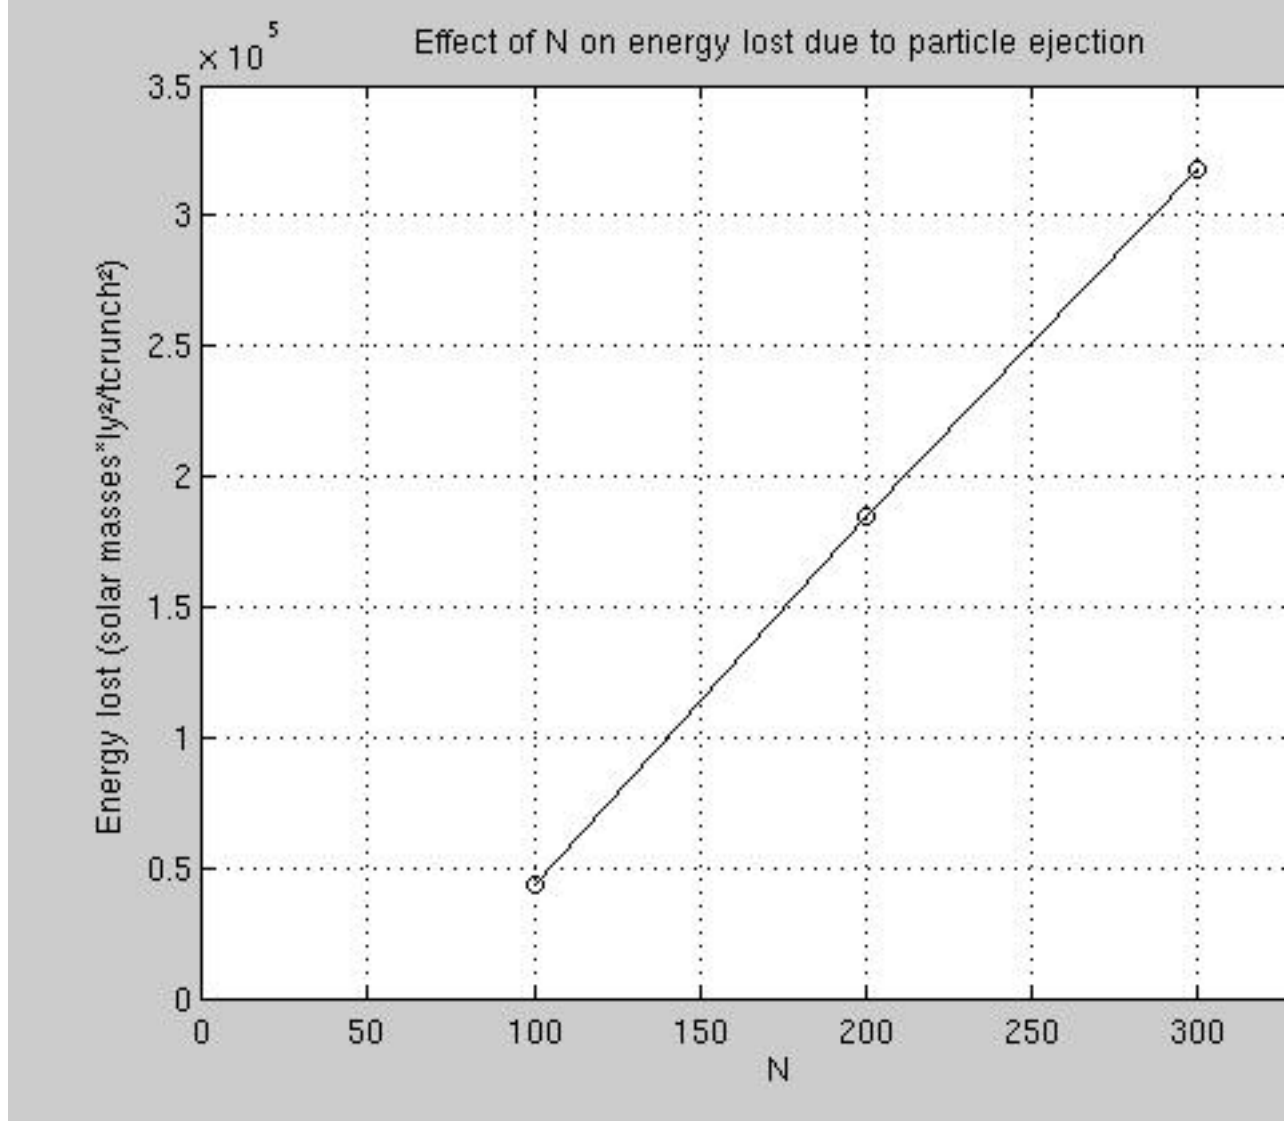


Figure 12: Lost energy as a function of N , using the ϵ scaling in (5).




Figure 13: Number of bound particles for N initial = 100 bodies, $\varepsilon = 0$

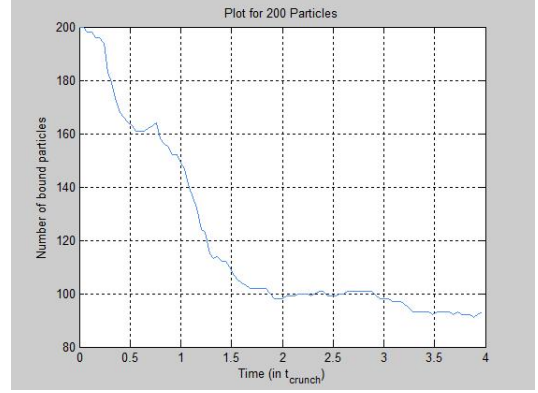


Figure 14: Number of bound particles for N initial = 200 bodies, $\varepsilon = 0$




Figure 15: Number of bound particles for N initial = 300 bodies, $\varepsilon = 0$

Without using the modified gravitational potential, the energy is clearly not conserved: a lot of particles are ejected out of the system. For these particles, the kinetic energy blows up, and the total energy of these particles is positive.




Figure 16: Number of bound particles for $N_{\text{initial}} = 100$, $\epsilon = 0.10005$




Figure 17: Number of bound particles for $N_{\text{initial}} = 200$, $\epsilon = 0.071$




Figure 18: Number of bound particles for $N_{\text{initial}} = 300$ bodies, $\epsilon = 0.06$

After adding the smoothing part in our gravitational potential, we can see that the quantity kinetic energy does not blow up anymore, and the quantity of energy lost is a lot smaller than before: this was expected, since we reduced the numerical instability by adding this ϵ factor in the calculation of our Newtonian force. Thus, the behavior of particles when they come closer to each other is more stable: we lose less particles than before.

The energy loss due to particle ejection also drops by a considerable amount when we increase ϵ , until $\epsilon \approx 0.4$, at which point random fluctuations take over. This is the same behaviour that was found for the fraction of bound particles with the same N .

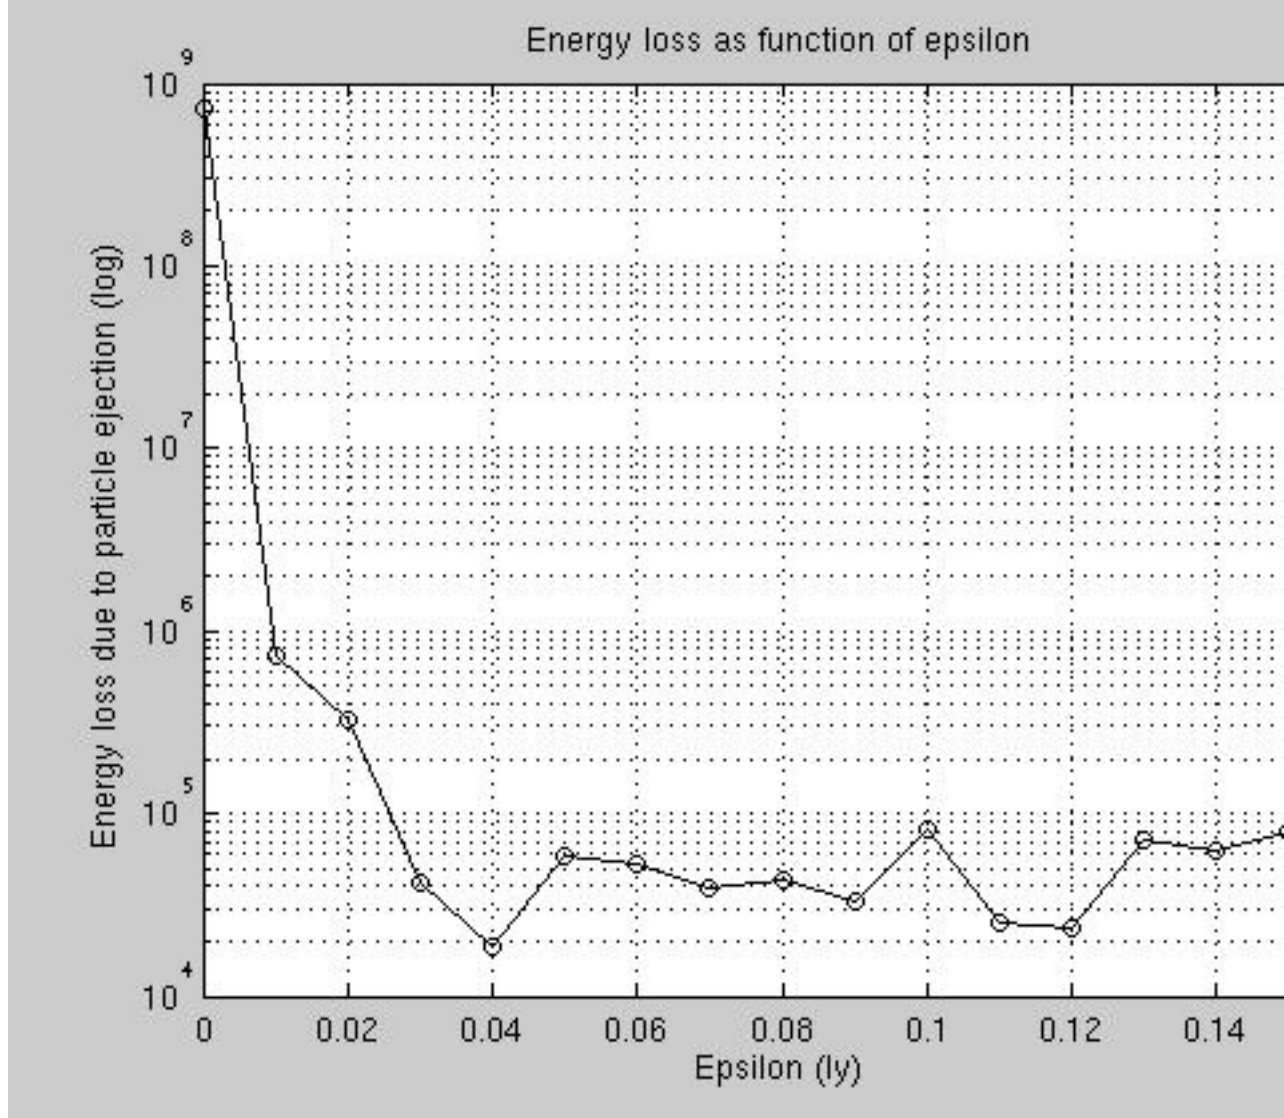


Figure 19: Energy loss ($\Delta E = \sum E_{all} - \sum E_{bound}$) for different values of ϵ , made using $N = 100$.

4.5 Virial Theorem




Figure 20: Ratio $\frac{\langle P \rangle}{\langle K \rangle}$ for different epsilons, $N = 100$

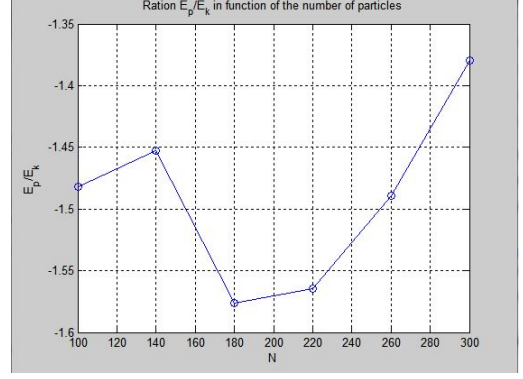


Figure 21: Ratio $\frac{\langle P \rangle}{\langle K \rangle}$ for different N , epsilon ≈ 0.03

After running the system for a few τ_{crunch} , we evaluate the kinetic and the potential energies for every particles of our bound system. What is displayed here is the value of the average kinetic energy, $\langle K \rangle$, and the average of the potential energy, $\langle P \rangle$, for our bound system. We can see that the virial theorem is met in the case where the smoothing function is used, with an ϵ of $\sqrt{0,0225}$, and is not met in our initial case (i.e. without using the smoothing function). The smoothing function helps us to take care of the numerical instability generated when two particles are too close. Indeed, with the standard gravitational potential, we lose too many particles. Thus, the introduction of ϵ solves this problem.

4.6 Radial distribution and density in equilibrium

The largest simulation we found it practical to make was a series of four simulations $N \in [500, 2000]$ using the auto-generated ϵ values. The total running time in this case was about 12.5 hours on a single processor core as expected for $\mathcal{O}(n^2)$ algorithms when the $N = 100$ case takes about a minute to run (depending on processor load and frequency).

For the largest simulation ($N = 2000$), we got the following radial distribution of particles:

In other words, we get a standard deviation that is about twice as large as the average. Looking at these values for all four simulations, we see that this is a general trend:

We attribute the large standard deviation to the fact that some particles, while still bound, have been excited to such high kinetic energy levels that they are

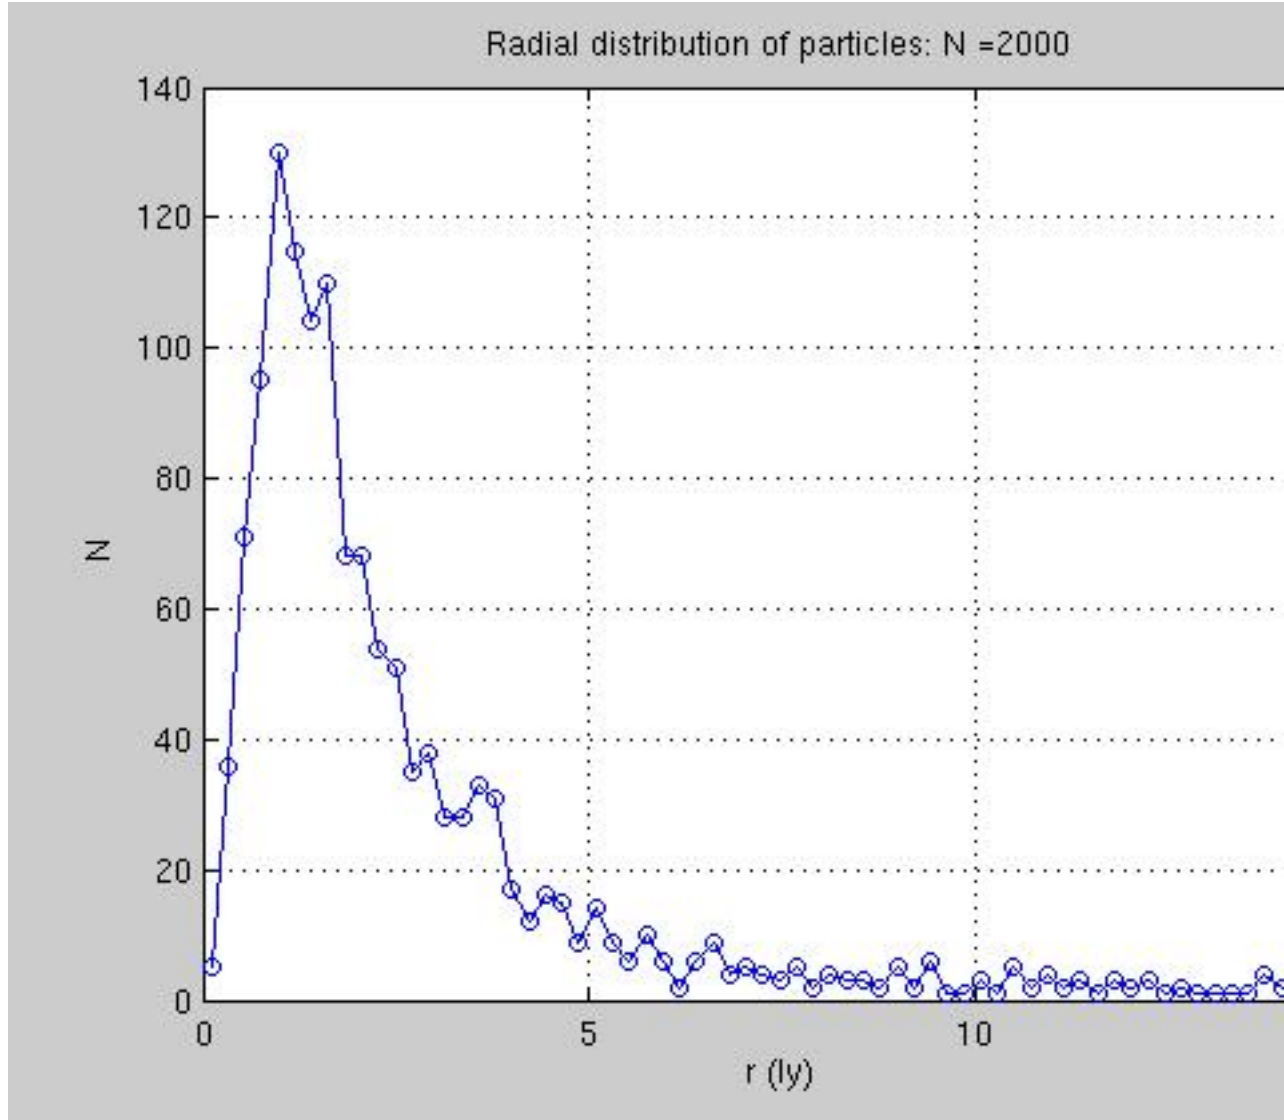


Figure 22: Radial distribution (histogram) for $N = 2000$, with r in the range $[0, \mu_{bound} + 1\sigma_{bound}]$. For the bound particles we got $\mu_{bound} = 4.83$ ly and $\sigma_{bound} = 9.68$ ly (measured from the bound particles' center of mass).

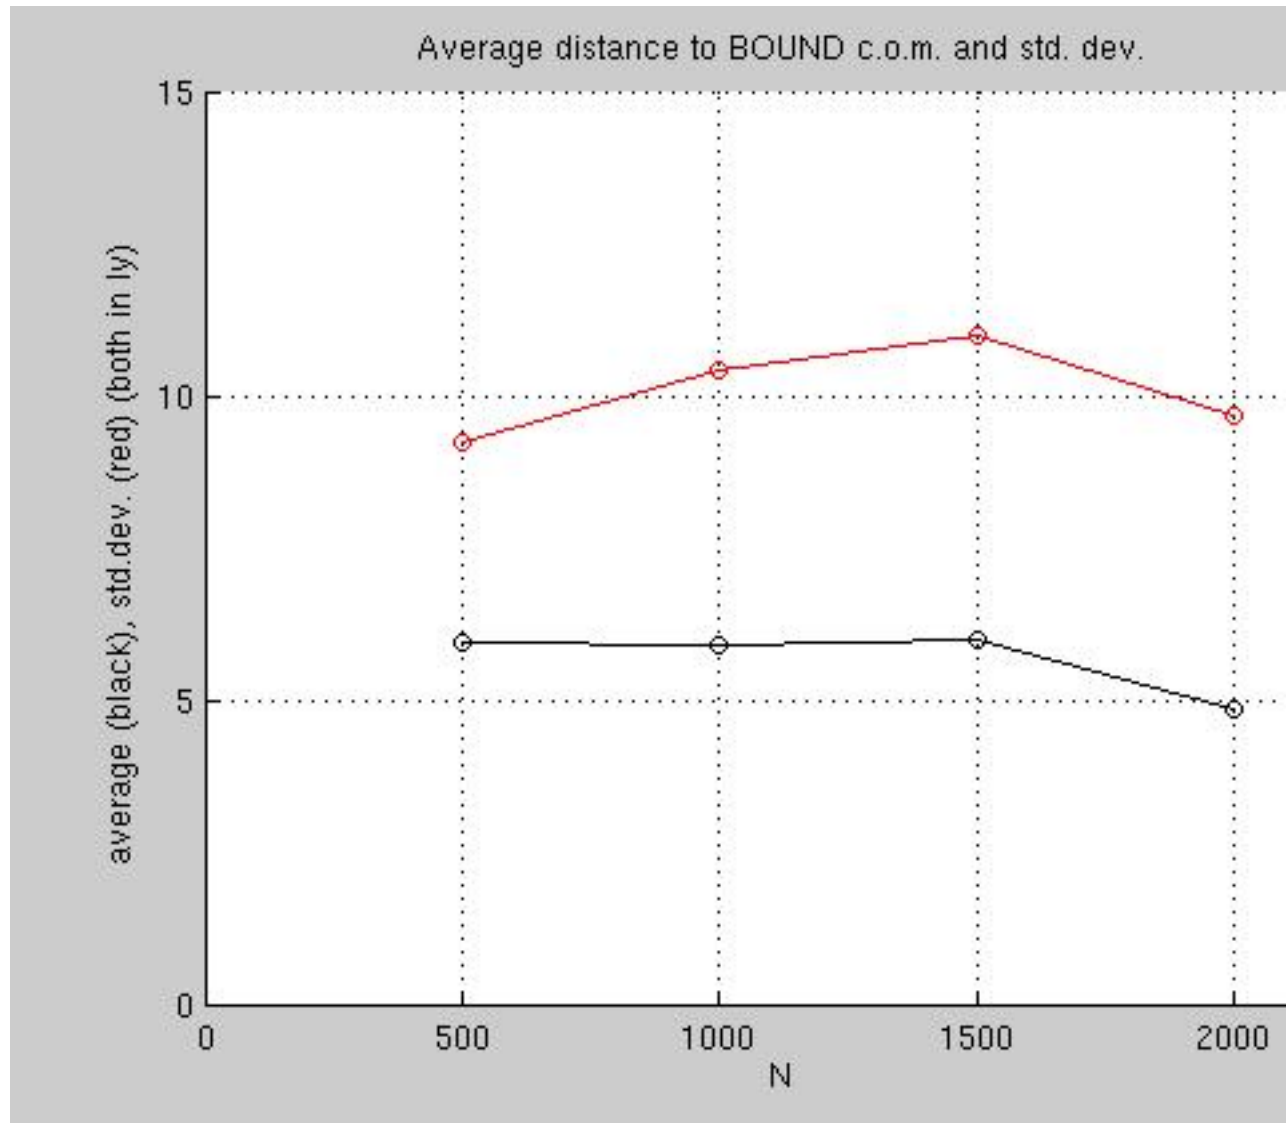


Figure 23: μ_{bound} (black) and σ_{bound} (red), plotted against N .

close to not being bound anymore (these would be the same particles responsible for the virial theorem not holding exactly even for the bound particles). Thus, we expect the standard deviation to vary roughly as the number of bound particles does - more bound particles means a higher change of high-energy bound particles (less energy lost to the unbound particles). It is clear, however, that we do not have sufficient data to find an expression for $\mu_{bound}(N)$, however we note that the radius does indeed shrink with increased N as noted in [4].

We then proceeded to plot the number density $n(r) = \frac{N(r)}{V(r)}$ for each spherical shell in our histogram and used the method of least squares to fit these data to a curve of the form

$$n(r) = \frac{n_0}{1 + \left(\frac{r}{r_0}\right)^4} \quad (9)$$

where we seek to determine the value of the constant $n_0 = n(0)$ (the core number density) and r_0 . The latter serves as a scale length for $n(r)$, as seen by the fact that $n(r_0) = \frac{n_0}{2}$, thus r_0 is the radius at which we find that the core density has initially halved (it is easily seen that this trend does not continue as we move further out, however).

We conclude that this gives a nice fit to our radial distribution.

5 Conclusion

x

Paralellizing the code would indeed be beneficial (due to time constraints, ironically, we were forced not to implement this as we would then need to find out how to do this in Visual C++). It is especially beneficial when doing series of many simulations that should result in a plottable data set over all the values of N we are interested in. One could then aim to distribute the load evenly over each processor. Furthermore, even with just one simulation, the $\mathcal{O}(n^2)$ parts of the code (gravitational forces, potential energy),

References

- [1] Ø. Elgarøy: Lecture notes in AST 4220, University of Oslo, fall 2009, pp. 31-33: <https://www.uio.no/studier/emner/matnat/astro/AST4220/h09/course-material/lectures.pdf>
- [2] Ø. Elgarøy: The Spherical Collapse Model: <https://www.uio.no/studier/emner/matnat/astro/AST4320/h12/undervisningsmateriale/spherecollapse.p>

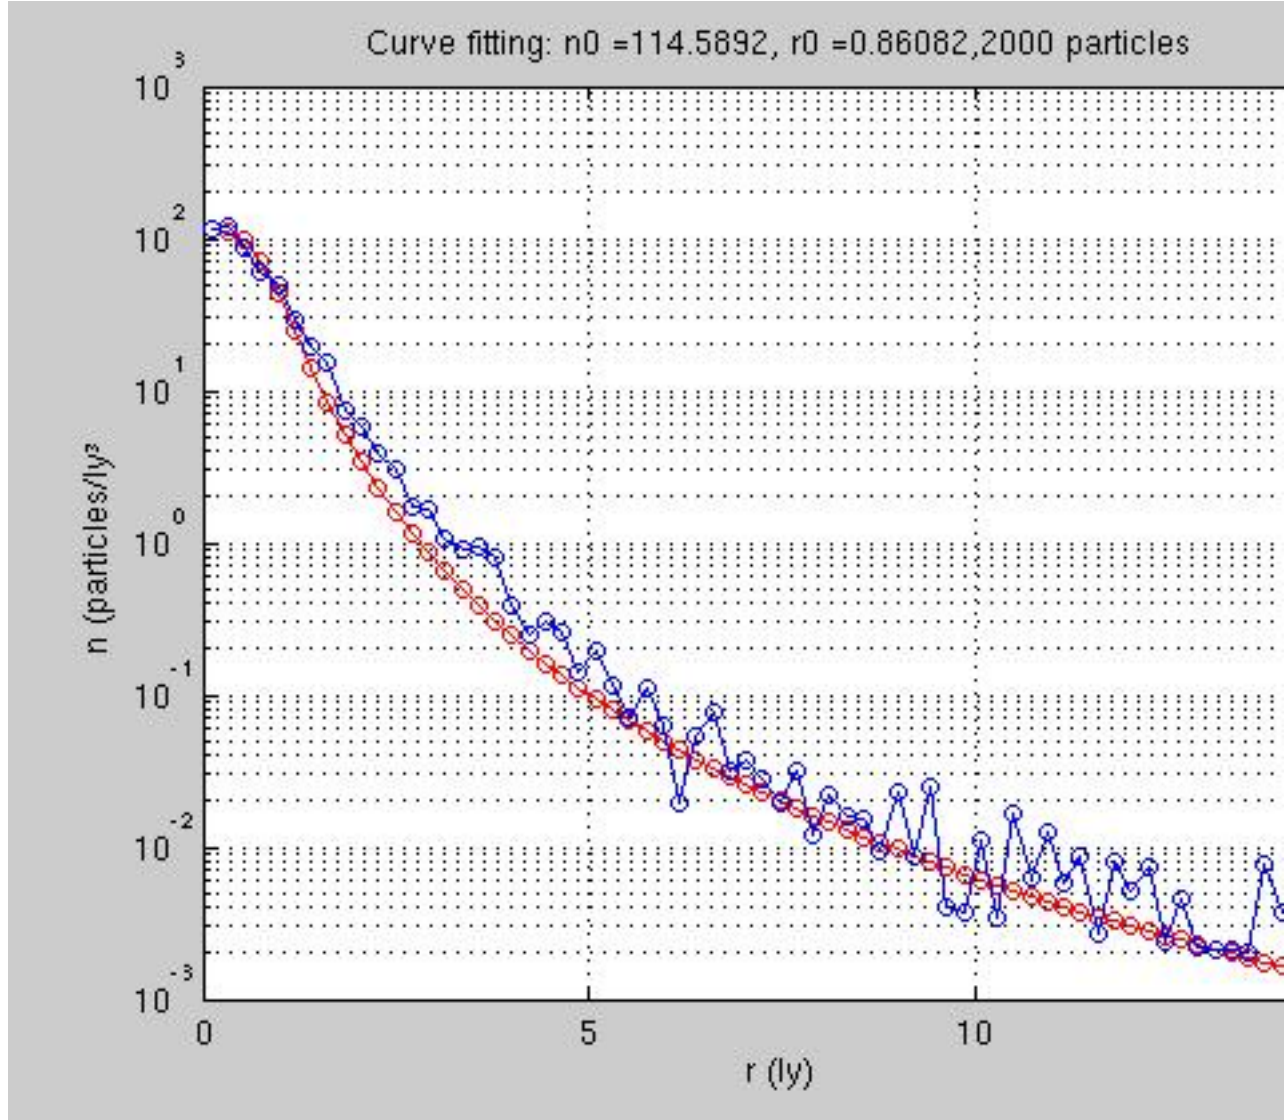


Figure 24: $n(r)$ (logarithmic) plotted against r for $N = 2000$. The blue curve is from our histogram, the red curve is given by (9) with $n_0 = 114.59 \text{ ly}^{-3}$, $r_0 = 0.86 \text{ ly}$.

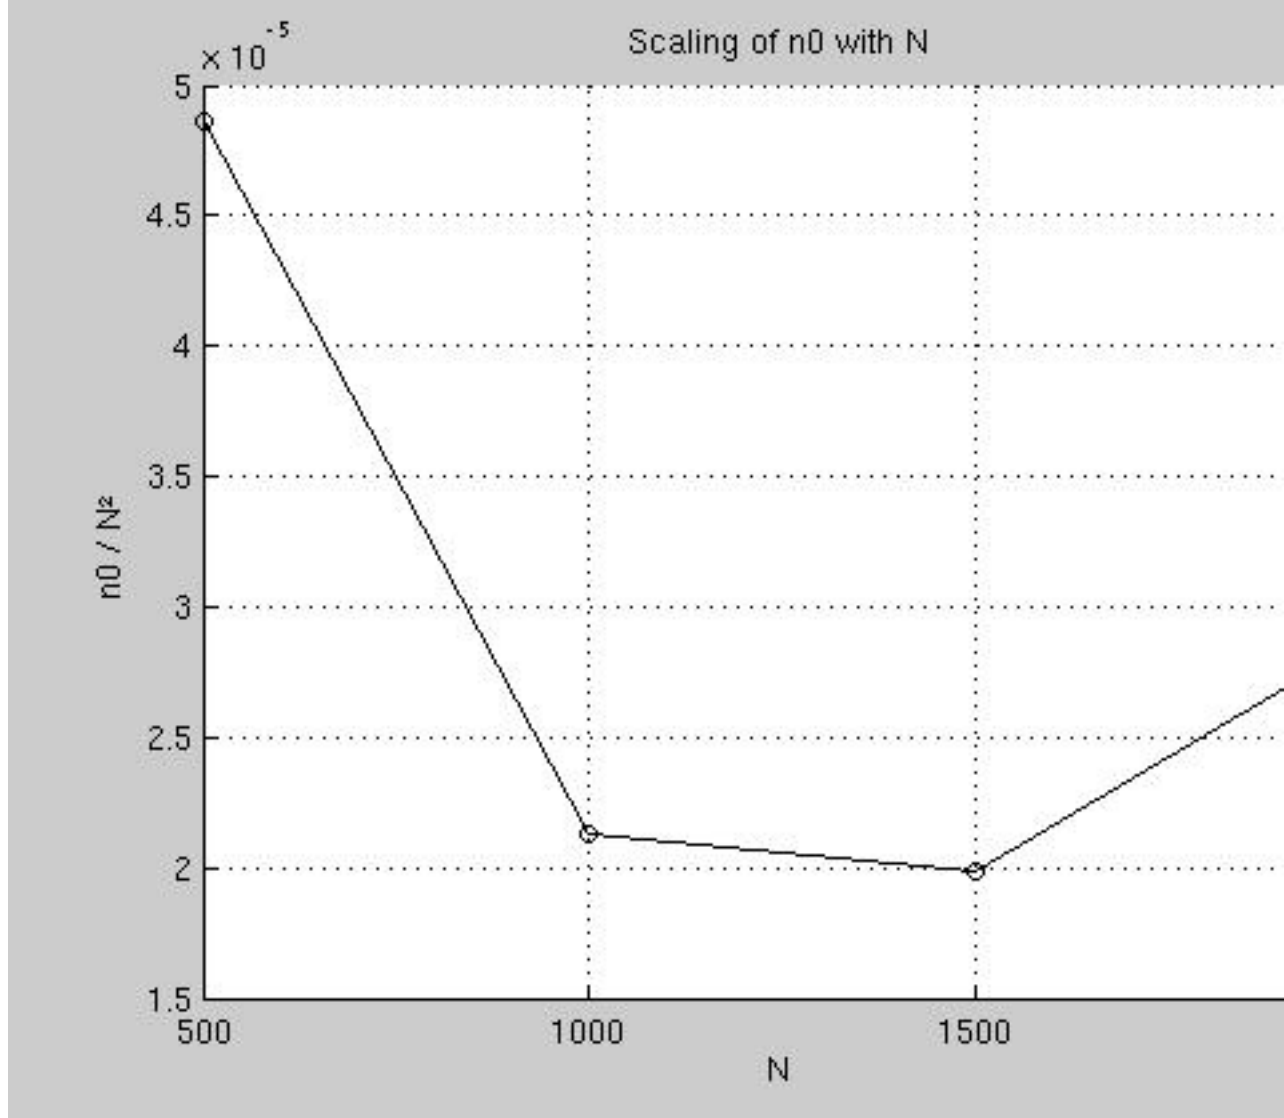


Figure 25: Comparison to the scaling behaviour found for $n(r)$ in [4]. We scaled the axes to match these predictions, but did not encounter the expected scaling behavior, where $n_0 \propto N^2$ and $r_0 \propto N^{-1/3}$: In that case, we would expect the values to remain more or less constant. However, it may well be that we are simply simulating too few particles to be able to reproduce this behavior. Another possible explanation is that we have not kept our ϵ values constant, but used the dynamic scaling of this parameter from (8).

- [3] J.A. Peacock: Large scale surveys and cosmic structure (Lectures delivered at the 2002 Tenerife Winter School, "Dark matter and dark energy in the universe"), pp. 15 - 19: <http://arxiv.org/abs/astro-ph/0309240>
- [4] M. Joyce, B. Marcos, F. Sylos Labini: "Cold uniform spherical collapse revisited", 2010: <http://arxiv.org/abs/1011.0614>
- [5] F.K. Hansen: "The Virial Theorem", lecture notes in AST 1100, University of Oslo, fall 2013: <https://www.uio.no/studier/emner/matnat/astro/AST1100/h13/undervisningsmateriale/lecture5.pdf>
- [6] R. Feres: "Math 350 Fall 2012 - Homework 8 Solutions", Washington University in St. Louis: <http://www.math.wustl.edu/~feres/Math350Fall2012/Math350F12HW08Sol.pdf>
- [7] Rubinstein: <title unknown>, photocopy from University of Colorado at Boulder: <http://ecee.colorado.edu/~fmeyer/class/ecen5322/rubinstein.pdf>
- [8] T. Clifton, P. G. Ferreira, A. Padilla, C. Skordis: "Modified Gravity and Cosmology", 2011: <http://arxiv.org/abs/1106.2476>